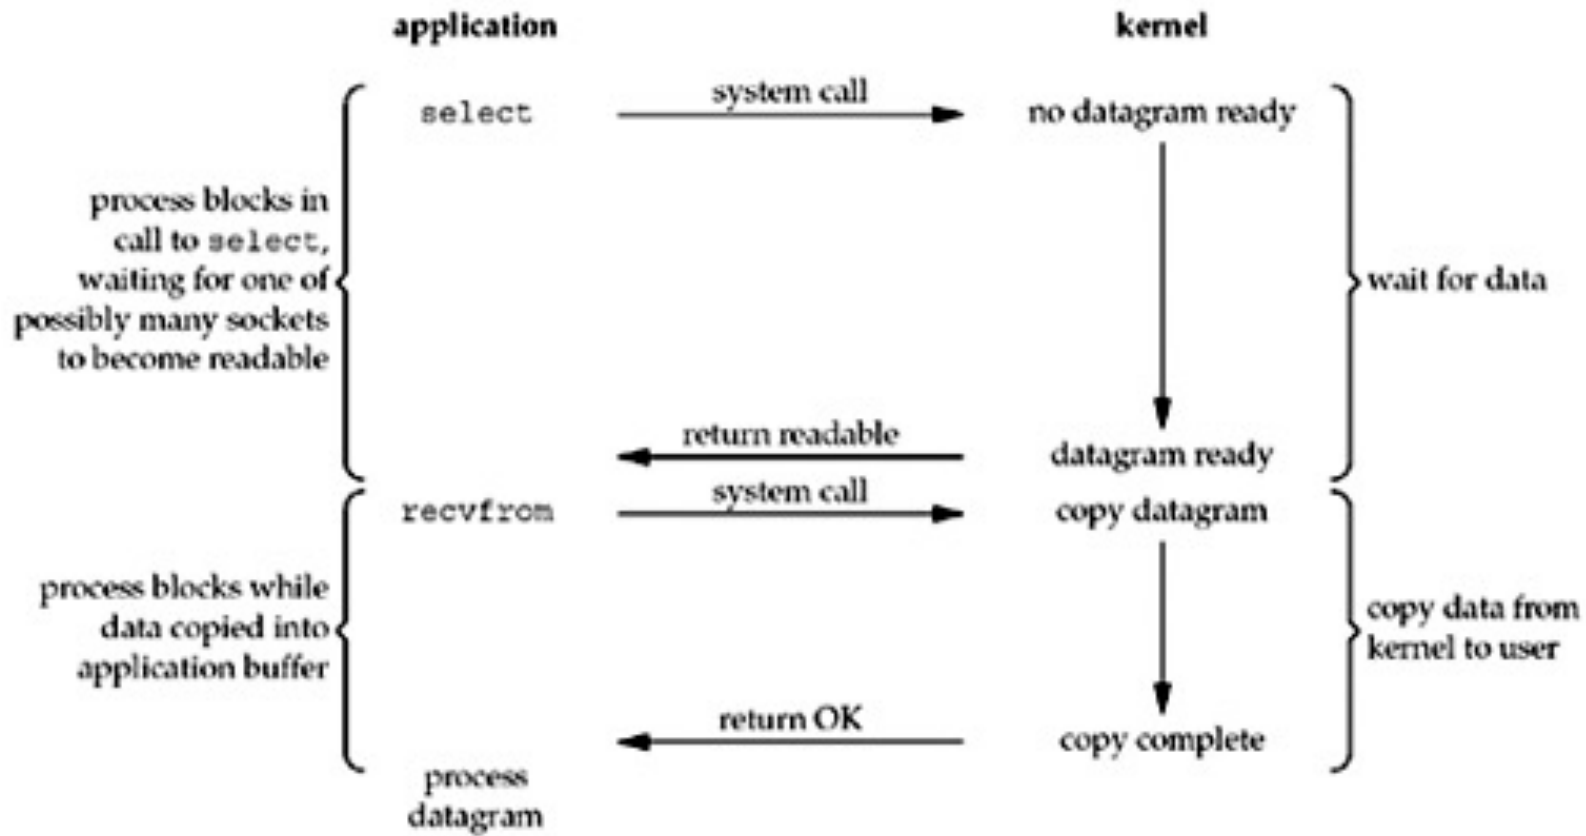


I/O MULTIPLEXING SERVER (CONT.)

Content

- I/O Multiplexing Model
- `select()`
- `poll()`

I/O Multiplexing Model



select()

- The **select()** function asks kernel to simultaneously check multiple sockets to see if they have data waiting to be **recv()**, or if you can **send()** data to them without blocking, or if some exception has occurred.
- The kernel to wake up the process only when one or more of events occurs or when a specified amount of time has passed.
- Exp : kernel to return only when
 - {1, 4, 5} are ready for reading
 - {2, 7} are ready for writing
 - {1, 4} have an exception condition pending
 - 10.2 seconds have elapsed

select()

```
#include <sys/select.h>
int select(int maxfd, fd_set *readfds, fd_set *writefds,
           fd_set *exceptfds, struct timeval *timeout);
```

- [IN] **maxfd** SHOULD BE the highest-numbered file descriptor in any of the three sets, plus 1
- [IN] **readfds**, **writefds**, **exceptfds** : set of FD to test for reading/receiving, writing/sending, exception conditions. At least one of them SHOULD BE not `NULL`.
- [IN] **timeout** : how long to wait for one of the specified descriptors to become ready. There are three types of using timeout
 - **NULL**: Wait forever
 - Wait up to a fixed amount of time
 - Do not wait at all : **timeout** points to the a `timeval` structure having the value 0
- Return:
 - On success, returns the total number of bits that are set(that is the number of ready file descriptors)
 - On time-out, returns 0
 - On error, return -1

fd_set

- Other *fd_sets* use to specify the descriptors that we want the kernel to test for reading, writing, and exception conditions.
- To specify one or more descriptor values for each of these three arguments, select uses *descriptor sets*.
- All the implementation details are irrelevant to the application and are hidden in the *fd_set* datatype and the following four macros:

```
void FD_ZERO(fd_set *fdset);      /* clear all bits in fdset */  
void FD_SET(int fd, fd_set *fdset); /* turn on the bit for fd in fdset */  
void FD_CLR(int fd, fd_set *fdset); /* turn off the bit for fd in fdset */  
int FD_ISSET(int fd, fd_set *fdset); /* Return true if fd is in the fdset */
```

select () - Conditions

- Ready for reading:
 - The number of bytes of data in the socket receive buffer is greater than or equal to the current size of the low-water mark for the socket receive buffer
 - The read half of the connection is closed
 - The listening socket receives a new connection request
 - A socket error is pending
- Ready for writing:
 - The size of the available space in the socket send buffer and either: (i) the socket is connected, or (ii) the socket does not require a connection (e.g., UDP).
 - The write half of the connection is closed
 - A socket using a non-blocking `connect` has completed the connection, or the `connect` has failed
 - A socket error is pending
- Exception: TCP out-of-band data

Examples

```
int s1, s2, n;
fd_set readfds;
struct timeval tv;
char buf1[256], buf2[256];
//pretend we've connected both to a server at this point
//s1 = socket(...); s2 = socket(...);
//connect(s1, ...)... connect(s2, ...)...

// clear the set ahead of time
FD_ZERO(&readfds);

// add our descriptors to the set
FD_SET(s1, &readfds);
FD_SET(s2, &readfds);

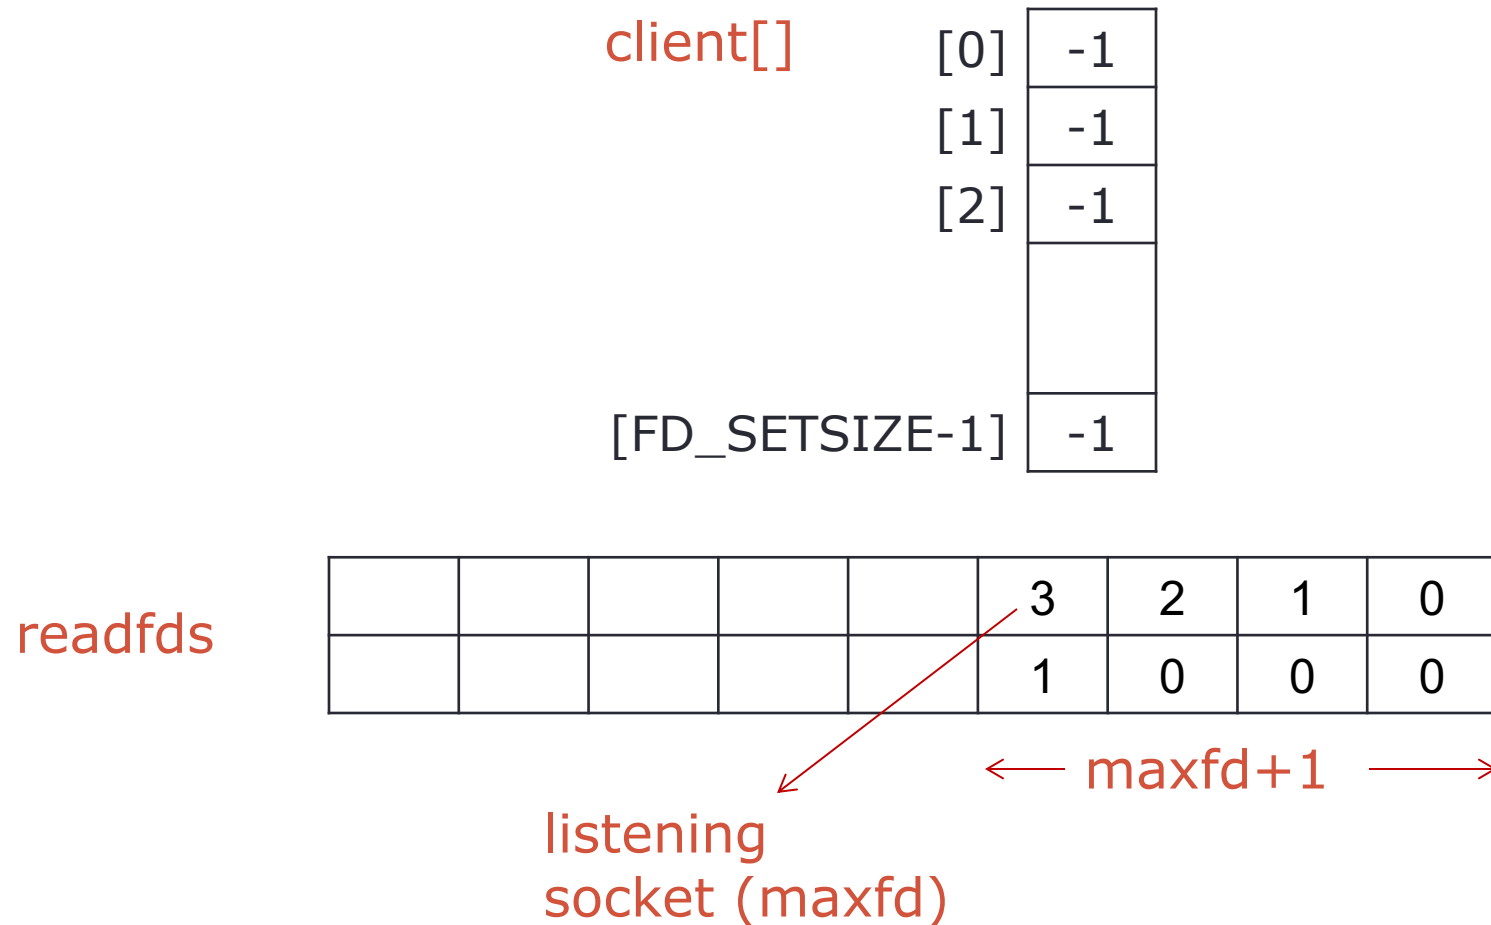
// since we got s2 second, it's the "greater", so we use
that for
// the n param in select()
n = s2 + 1;
```


Examples (2)

```
// wait until either socket has data ready to be recv()d
//(timeout 10.5 secs)
tv.tv_sec = 10;
tv.tv_usec = 500000;
rv = select(n, &readfds, NULL, NULL, &tv);
if (rv == -1) {
    perror("\Error: "); // error occurred in select() }
else if (rv == 0)
    printf("Timeout occurred! No data after 10.5s \n");
else {
    // one or both of the descriptors have data
    if (FD_ISSET(s1, &readfds))
        recv(s1, buf1, sizeof buf1, 0);
    if (FD_ISSET(s2, &readfds))
        recv(s1, buf2, sizeof buf2, 0);
}
```

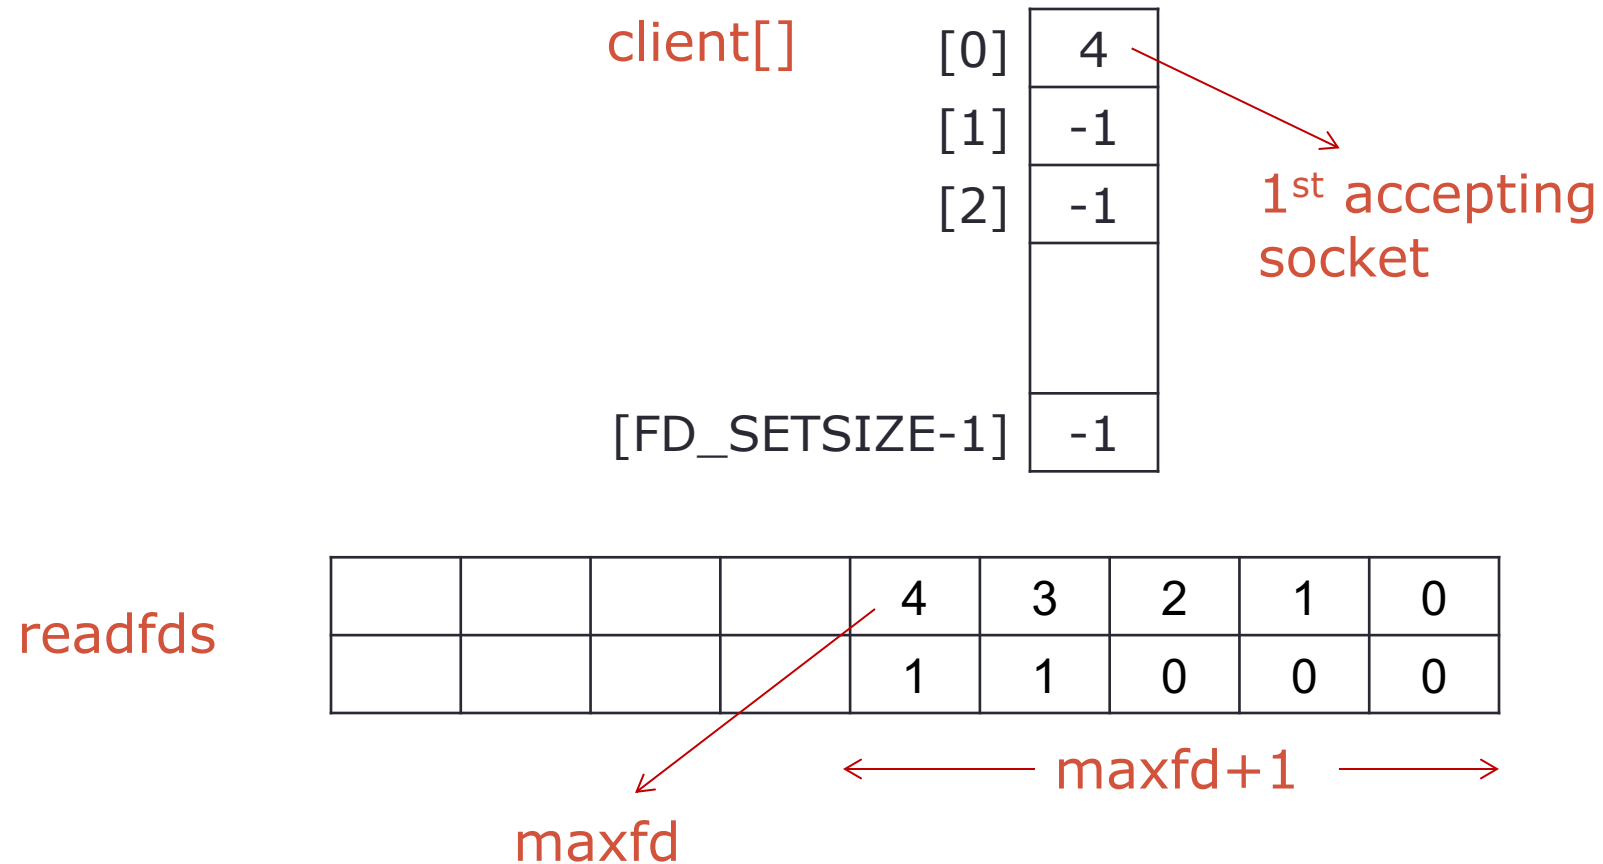
How to use select() in TCP server

Data structures for TCP server with just a listening socket



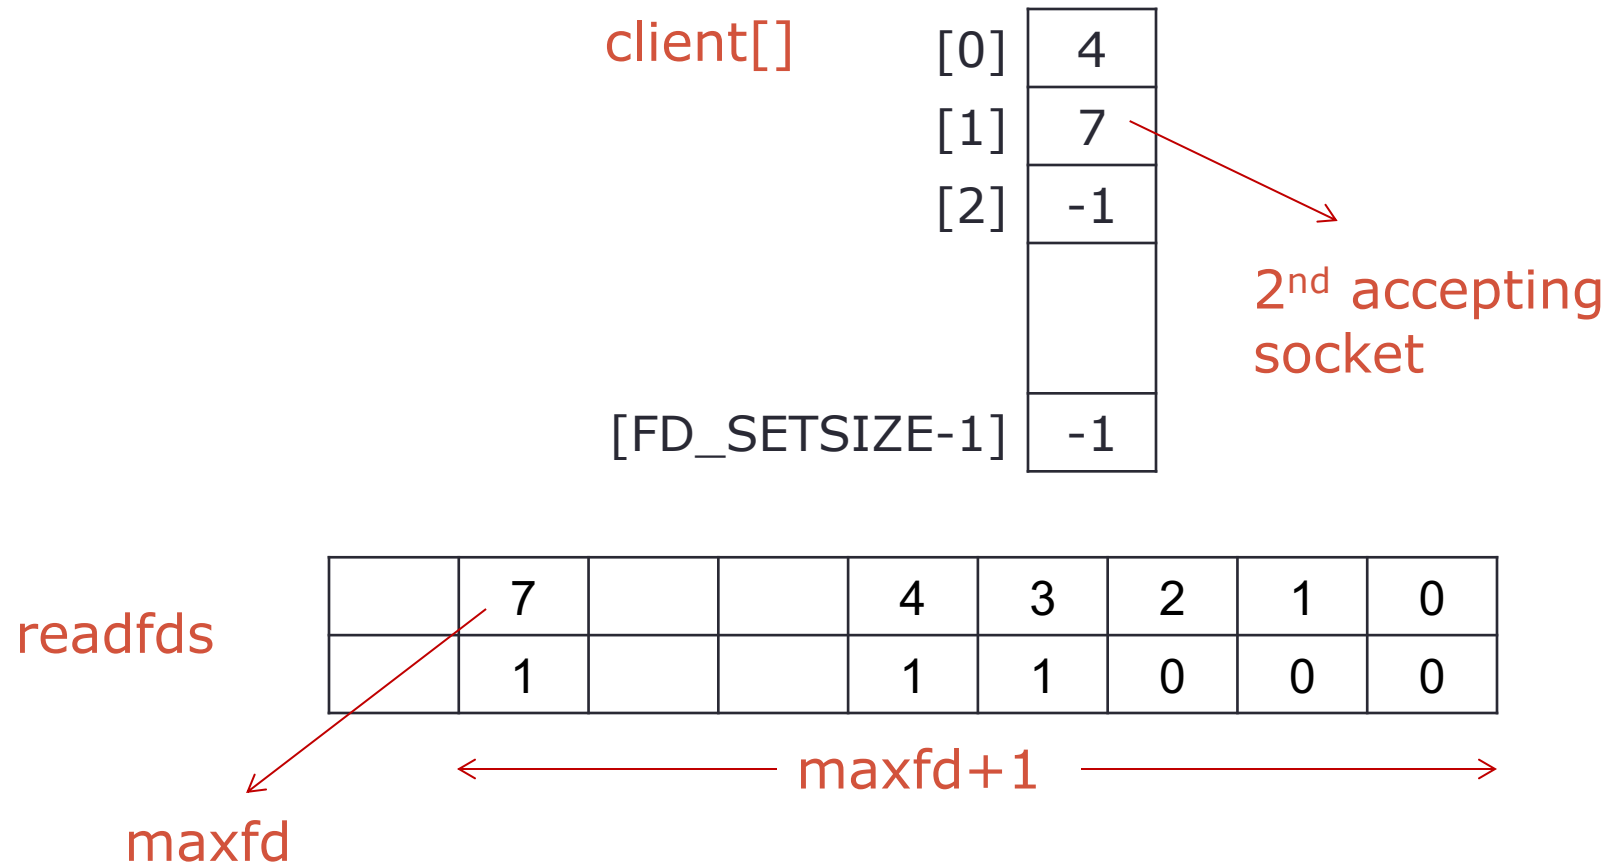
How to use select() in TCP server

Data structures after the 1st client connection is established



How to use select() in TCP server

Data structures after the 2nd client connection is established



How to use select() in TCP server

After `select()` return. Example:

- New connection has established
- No data arrival on socket 7
- Socket 4 is ready for reading

`client[]`

[0]	4
[1]	7
[2]	-1
[FD_SETSIZE-1]	-1

2nd accepting socket

`readfds`

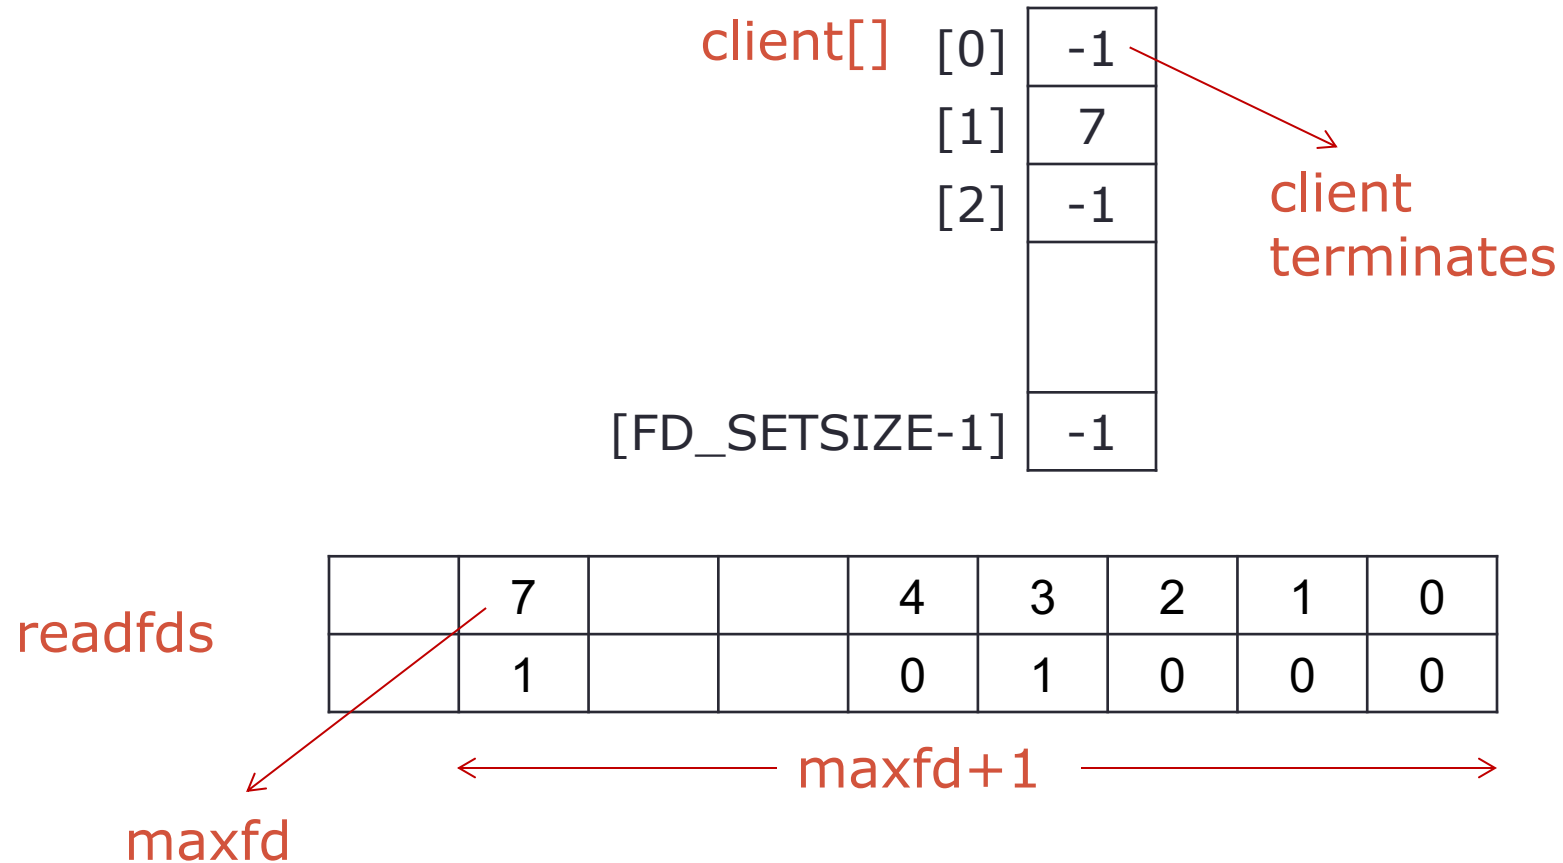
	7			4	3	2	1	0
	0			1	1	0	0	0

maxfd

maxfd + 1

How to use select() in TCP server

Data structures after a client terminates its connection



How to use select() in TCP server

```
listenfd = socket(...);  
listen(listenfd, ...);  
maxfd = listenfd;
```

```
//Assign initial value for the array of connection socket  
for(...) client[i] = -1;
```

```
//Assign initial value for the fd_set  
FD_ZERO (...);
```

```
//Set bit for listenfd  
FD_SET(listenfd, ...)
```

How to use select() in TCP server

```
//Communicate with clients
while(...){
    nEvents = select(...);
    //check the status of listenfd
    if(FD_ISSET(listenfd,...)){
        connfd = accept(...);
        maxfd = connfd;
        if (client[i] == -1)
            client[i] = connfd;
    }
    //check the status of connfd(s)
    for(...){
        if(FD_ISSET(client[i],...)){
            doSomething();
            close(connfd);
            client[i] = -1;
            FD_CLEAR(client[i],...)
        }
    }
}
```